
AlphaDou: Multi-Agent Search and Learning for Jungle Chess

Zihan Wang
wangzh12023

Ziyu Cheng
chengzy2023

Mingfei Xia
xiamf2023

Abstract

In a simplified Jungle Chess environment (7×8 board, omitting traps, rivers, and dens), we implemented and compared seven agents: Random, Greedy (heuristic), Minimax, Tabular Q-Learning, Approximate QLearning and Deep Q-Network (DQN). Each agent was evaluated via self-play and round-robin tournaments, comparing win rates and decision-making latency. We also conducted ablation studies to assess the impact of heuristic weights and the discount factor on learning performance. Our results demonstrate the applicability of various search and learning methods to Jungle Chess and provide a foundation for future reinforcement learning research incorporating neural priors or full game rules.

1 Introduction

1.1 Game Settings

- **Board:** 7×8 Grid
- **Pieces:** 8 pieces per side: Elephant (Rank 8), Lion (Rank 7), Tiger (Rank 6), Leopard (Rank 5), Wolf (Rank 4), Dog (Rank 3), Cat (Rank 2), Rat (Rank 1)
- **Simplified Rules:**
 - Rankbased Predation: A piece of higher rank can capture (eat) any opposing piece of strictly lower rank. But a Rat (Rank 1) can capture an Elephant (Rank 8)
 - No Traps / No Rivers / No Dens:
 - Action: Each turn, a player must either move one revealed piece to an adjacent (up/down/left/right) empty square or reveal one facedown pieces if any remain.
 - Game End Condition:
 - * Win by Elimination: A player who captures all of the opponents remaining pieces immediately wins.
 - * Stalemate Rule: If no capture occurs in 100 consecutive moves (both sides combined), the game is declared a draw.

1.2 State Space

1.2.1 State Representation

- A *state* is fully determined by the set of all remaining pieces and their board positions.
- In our simplified rules, there are no traps, rivers, or dens, so all 16 pieces (8 per side) occupy one of the 56 squares.
- We encode a state as an ordered tuple of 16 triples:

$$s = ((p_1, x_1, y_1), (p_2, x_2, y_2), \dots, (p_{16}, x_{16}, y_{16})),$$

where:




Figure 1: Visualization of Jungle Chess game board

- $p_i \in \{\text{Black Elephant}, \dots, \text{White Rat}\}$ specifies piece identity (type and color), and
- (x_i, y_i) is the square coordinate (row and column) of that piece.
- Once two pieces occupy the same square, the lowerrank piece is captured (removed), so no square holds more than one piece.

1.2.2 Combinatorial Size

- There are 56 total squares and 16 distinct pieces.
- To place 16 Black pieces on 56 squares and order matters: $\binom{56}{16} \times 16!$
- And every piece has two states: revealed or unrevealed: $\binom{56}{16} \times 16! \times 2^{16}$
- And two people's turn counts as two different situations: $\binom{56}{16} \times 16! \times 2^{16} \times 2$

$$|\mathcal{S}| = \binom{56}{16} \times 16! \times 2^{16} \times 2 \approx 2^{32}$$

1.2.3 Action Space

- In any given state s , a players legal actions consist of:
 - Moving one revealed piece to an adjacent empty square (up, down, left, or right), provided it does not violate rankbased capture rules.
 - Flipping (revealing) one of their own facedown pieces, if any remain.
- Each revealed piece can have up to 4 orthogonal moves, subject to:
 - The destination square must be empty or occupied by an opponents piece of strictly lower rank (with the RatElephant exception).
 - No diagonal moves are allowed.

2 Methods

2.1 RandomAgent

The Random Agent implements the simplest decision-making strategy in our Chinese Dark Chess system. It serves as a baseline player that makes completely random moves without any strategic consideration or game state evaluation. The Random Agent follows this simple procedure:

1. **Action Retrieval:** Obtains all possible legal actions from the current board state.
2. **Randomization:** Shuffles the list of possible actions using.
3. **Action Execution:** Iterates through the shuffled actions and attempts to execute the first valid one:

2.2 GreedyAgent

The Greedy Agent implements a heuristic-based decision-making strategy that evaluates all possible moves and selects the action that provides the highest immediate score improvement. Unlike the Random Agent, it incorporates strategic thinking through a comprehensive evaluation function, but without lookahead planning like more sophisticated algorithms.

The GreedyPlayer class uses a multi-component evaluation function based on the comprehensive formula:

$$\text{Score}(s) = \sum_{i=1}^4 \text{Component}_i(s)$$

Where the components and parameters are defined as follows:

1. Piece Value Component:

This component evaluates the material advantage of the agent.

$$\text{Component}_1(s) = \sum_{p \in P_{\text{self}}} V(p) - \lambda_1 \sum_{p \in P_{\text{opp}}} V(p) \quad (1)$$

Here:

- P_{self} and P_{opp} are the sets of pieces belonging to the self and the opponent respectively.
- $V(p)$ is the value function for a piece p :

$$V(p) = \begin{cases} \text{piece_values}[\text{strength}(p)] & \text{if } p.\text{revealed} = \text{True} \\ V_{\text{unrevealed}} & \text{if } p.\text{revealed} = \text{False} \end{cases} \quad (2)$$

- piece_values is a mapping from piece strength to its intrinsic value:

$$\text{piece_values} = \{8 \mapsto 80, \quad 7 \mapsto 70, \quad 6 \mapsto 60, \quad 5 \mapsto 50, \\ 4 \mapsto 40, \quad 3 \mapsto 30, \quad 2 \mapsto 20, \quad 1 \mapsto 30\}$$

- $V_{\text{unrevealed}} = 30$ is the estimated value for an unrevealed piece.
- $\lambda_1 = 0.8$ is the opponent piece weight coefficient.

2. Positional Reward Component:

This component rewards pieces for occupying advantageous board positions.

$$\text{Component}_2(s) = \lambda'_3 \sum_{p \in P_{\text{self}}^{\text{revealed}}} (f_{\text{advance}}(p) + f_{\text{center}}(p)) \quad (3)$$

Here:

- $P_{\text{self}}^{\text{revealed}}$ is the set of revealed pieces belonging to the agent.
- $f_{\text{advance}}(p)$ is the advance reward function for piece p :

$$f_{\text{advance}}(p) = \begin{cases} 5 \times \text{row}(p) & \text{if } \text{player_id} = 0 \\ 5 \times (\text{ROWS} - 1 - \text{row}(p)) & \text{if } \text{player_id} = 1 \end{cases} \quad (4)$$

where $\text{row}(p)$ is the row index of piece p , ROWS is the total number of rows on the board, and player_id identifies the current player.

- $f_{\text{center}}(p)$ is the central control reward function for piece p :

$$f_{\text{center}}(p) = \begin{cases} 5 & \text{if } \text{col}(p) \in \{\lfloor \text{COLS}/2 \rfloor - 1, \lfloor \text{COLS}/2 \rfloor\} \\ 0 & \text{otherwise} \end{cases} \quad (5)$$

where $\text{col}(p)$ is the column index of piece p , and COLS is the total number of columns on the board.

- $\lambda'_3 = \frac{\lambda_3}{10} = \frac{15}{10} = 1.5$ is the scaled positional weight coefficient. (Original $\lambda_3 = 15$).

3. Mobility Component:

This component rewards the agent for having a larger number of available moves.

$$\text{Component}_3(s) = \lambda_2 |A_{\text{possible}}(s)| \quad (6)$$

Here:

- $A_{\text{possible}}(s)$ is the set of all legal actions available to the agent in state s .
- $|A_{\text{possible}}(s)|$ denotes the number of possible actions.
- $\lambda_2 = 2$ is the mobility weight coefficient.

4. Mobility Component:

This component evaluates the immediate threats and advantages in engagements between adjacent pieces.

$$\text{Component}_4(s) = w_{\text{safety}} \sum_{p_s \in P_{\text{self}}^{\text{revealed}}} \sum_{p_o \in P_{\text{opp}}^{\text{revealed}}} \text{ThreatValue}(p_s, p_o) \quad (7)$$

Here:

- $P_{\text{self}}^{\text{revealed}}$ and $P_{\text{opp}}^{\text{revealed}}$ are the sets of revealed pieces for the agent and opponent, respectively.
- $\text{ThreatValue}(p_s, p_o)$ is the function evaluating the threat between an agent's piece p_s and an opponent's piece p_o :

$$\text{ThreatValue}(p_s, p_o) = \begin{cases} -10 & \text{if } \text{Adjacent}(p_s, p_o) \wedge \text{CanCapture}(p_o, p_s) \\ +15 & \text{if } \text{Adjacent}(p_s, p_o) \wedge \text{CanCapture}(p_s, p_o) \\ 0 & \text{otherwise} \end{cases} \quad (8)$$

- $\text{Adjacent}(p_1, p_2)$ is a boolean function that returns true if pieces p_1 and p_2 are adjacent on the board.
- $\text{CanCapture}(p_1, p_2)$ is a boolean function determining if piece p_1 can capture piece p_2 :

$$\text{CanCapture}(p_1, p_2) = \begin{cases} \text{True} & \text{if } (\text{strength}(p_1) > \text{strength}(p_2)) \\ & \wedge \neg(\text{strength}(p_1) = 8 \wedge \text{strength}(p_2) = 1) \\ \text{True} & \text{if } \text{strength}(p_1) = 1 \wedge \text{strength}(p_2) = 8 \\ \text{False} & \text{otherwise} \end{cases} \quad (9)$$

where $\text{strength}(p)$ denotes the strength of piece p .

- $w_{\text{safety}} = 10$ is the safety weight coefficient.

The overall heuristic $H(s)$ combines these components using their respective weights: $\lambda_1 = 0.8$, $\lambda_2 = 2$, $\lambda'_3 = 1.5$, and $w_{\text{safety}} = 10$.

The heuristic evaluation function, $H(s)$, for the GreedyAgent is a weighted sum of four components: piece value, positional reward, mobility, and safety. The state of the game is denoted by s .

$$\begin{aligned} H(s) = & \left(\sum_{p \in P_{\text{self}}} V(p) - \lambda_1 \sum_{p \in P_{\text{opp}}} V(p) \right) \\ & + \lambda'_3 \sum_{p \in P_{\text{self}}^{\text{revealed}}} (f_{\text{advance}}(p) + f_{\text{center}}(p)) \\ & + \lambda_2 |A_{\text{possible}}(s)| \\ & + w_{\text{safety}} \sum_{p_s \in P_{\text{self}}^{\text{revealed}}} \sum_{p_o \in P_{\text{opp}}^{\text{revealed}}} \text{ThreatValue}(p_s, p_o) \end{aligned} \quad (10)$$

2.3 MinimaxAgent

The evaluation function for the MinimaxAgent, denoted as $H_{\text{minimax}}(s)$, assesses the desirability of a game state s . It is composed of four distinct components, each contributing to the overall score.

The complete Minimax evaluation function is defined as:

$$\begin{aligned}
H_{\text{minimax}}(s) = & w_1(|P_{\text{self}}| - |P_{\text{opp}}|) \\
& + w_2 \left(\sum_{p \in P_{\text{self}}^{\text{revealed}}} \text{strength}(p) - \sum_{p \in P_{\text{opp}}^{\text{revealed}}} \text{strength}(p) \right) \\
& + (w_{3a}|P_{\text{self}}^{\text{unrevealed}}| - w_{3b}|P_{\text{opp}}^{\text{unrevealed}}|) \\
& + \sum_{p_s \in P_{\text{self}}^{\text{revealed}}} \sum_{p_o \in P_{\text{opp}}^{\text{revealed}}} \text{TacticalValue}(p_s, p_o)
\end{aligned} \tag{11}$$

where the specific weights and components are detailed below:

1. Piece Count Difference Component

This component measures the raw numerical advantage in pieces.

$$C_1(s) = w_1(|P_{\text{self}}| - |P_{\text{opp}}|) \tag{12}$$

Where:

- $|P_{\text{self}}|$ is the total number of pieces for the agent (self).
- $|P_{\text{opp}}|$ is the total number of pieces for the opponent.
- $w_1 = 1$.

2. Piece Strength Sum Component

This component evaluates the difference in the sum of strengths of revealed pieces.

$$C_2(s) = w_2 \left(\sum_{p \in P_{\text{self}}^{\text{revealed}}} \text{strength}(p) - \sum_{p \in P_{\text{opp}}^{\text{revealed}}} \text{strength}(p) \right) \tag{13}$$

Where:

- $P_{\text{self}}^{\text{revealed}}$ and $P_{\text{opp}}^{\text{revealed}}$ are the sets of revealed pieces for the self and opponent respectively.
- $\text{strength}(p) \in \{1, 2, 3, 4, 5, 6, 7, 8\}$ denotes the strength of piece p .
- $w_2 = 10$

3. Information Value Component

This component quantifies the value of information asymmetry based on unrevealed pieces.

$$C_3(s) = w_{3a}|P_{\text{self}}^{\text{unrevealed}}| - w_{3b}|P_{\text{opp}}^{\text{unrevealed}}| \tag{14}$$

Where:

- $P_{\text{self}}^{\text{unrevealed}}$ and $P_{\text{opp}}^{\text{unrevealed}}$ are the sets of unrevealed pieces for the agent and opponent, respectively.
- $w_{3a} = 5$ is the value assigned to each of the agent's unrevealed pieces (potential resources).
- $w_{3b} = 10$ is the negative value assigned to each of the opponent's unrevealed pieces (unknown threats).

This component is used directly in $H_{\text{minimax}}(s)$ as defined.

4. Threat Opportunity Component

This component assesses immediate tactical threats and opportunities between adjacent revealed pieces.

$$C_4(s) = \sum_{p_s \in P_{\text{self}}^{\text{revealed}}} \sum_{p_o \in P_{\text{opp}}^{\text{revealed}}} \text{TacticalValue}(p_s, p_o) \quad (15)$$

Where:

- p_s is a revealed piece belonging to the agent, and p_o is a revealed piece belonging to the opponent.
- $\text{TacticalValue}(p_s, p_o)$ is the function evaluating the tactical situation:

$$\text{TacticalValue}(p_s, p_o) = \begin{cases} +20 & \text{if } \text{Adjacent}(p_s, p_o) \wedge \text{CanCapture}(p_s, p_o) \\ -15 & \text{if } \text{Adjacent}(p_s, p_o) \wedge \text{CanCapture}(p_o, p_s) \\ 0 & \text{otherwise} \end{cases} \quad (16)$$

- $\text{Adjacent}(p_1, p_2)$ is a boolean function, true if pieces p_1 and p_2 are adjacent.
- $\text{CanCapture}(p_1, p_2)$ is a boolean function indicating if p_1 can capture p_2

2.4 Q-Learning

2.4.1 Basic State Definition

The state tensor $\mathbf{S} \in \mathbb{R}^{7 \times 8 \times 4}$ is defined where each position (r, c) has a feature vector:

$$\mathbf{s}_{r,c} = \begin{bmatrix} p_{r,c} \\ \frac{\text{strength}_{r,c}}{8} \\ \mathbb{I}_{\text{revealed}}(r, c) \\ \mathbb{I}_{\text{occupied}}(r, c) \end{bmatrix} \in \mathbb{R}^4$$

Where:

$$p_{r,c} = \begin{cases} 0 & \text{if piece at } (r, c) \text{ belongs to player 0} \\ 1 & \text{if piece at } (r, c) \text{ belongs to player 1} \\ 0 & \text{if position } (r, c) \text{ is empty} \end{cases}$$

$$\frac{\text{strength}_{r,c}}{8} = \begin{cases} \frac{k}{8} & \text{if piece at } (r, c) \text{ has strength } k \in \{1, 2, \dots, 8\} \\ 0 & \text{if position } (r, c) \text{ is empty} \end{cases}$$

$$\mathbb{I}_{\text{revealed}}(r, c) = \begin{cases} 1 & \text{if piece at } (r, c) \text{ is revealed} \\ 0 & \text{otherwise} \end{cases}$$

$$\mathbb{I}_{\text{occupied}}(r, c) = \begin{cases} 1 & \text{if position } (r, c) \text{ contains a piece} \\ 0 & \text{if position } (r, c) \text{ is empty} \end{cases}$$

And the **Action Space** is defined as

- **Reveal Action:** ("reveal", (row, col)) for unrevealed own pieces
- **Move Action:** ("move", (start_row, start_col), (end_row, end_col))

2.4.2 Learning process

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t) \right]$$

Where the **Hyperparameters** are: $\alpha = 0.1, \gamma = 0.95, \varepsilon$ -greedy \downarrow

$$r_{t+1} = \text{REWARDFUNCTION}$$

$$\delta_t = r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t)$$

$$Q_{\text{new}}(s_t, a_t) = Q(s_t, a_t) + \alpha \cdot \delta_t$$

2.4.3 RewardFunction

The REWARDFUNCTION is defined as below:

$$r_{t+1} = R_{\text{Pos}}(s_t, a_t) + R_{\text{terminal}}(s_{t+1})$$

$$V(k) = \begin{cases} 3.0 & \text{if } k = 1 \text{ (Rat)} \\ 1.0 & \text{if } k = 2 \\ 1.5 & \text{if } k = 3 \\ 2.0 & \text{if } k = 4 \\ 2.5 & \text{if } k = 5 \\ 3.0 & \text{if } k = 6 \\ 3.5 & \text{if } k = 7 \\ 4.0 & \text{if } k = 8 \text{ (Elephant)} \end{cases}$$

$$R_{\text{Pos}}(s_t, a_t) = \frac{1}{n} \sum_i P(i) \times V(i) \times (\Delta(\text{Opp}) - \Delta(\text{Threat}))$$

Where the $P(i)$ is the probability of the i -th unrevealed piece to be at the selected position.

$$P(i) = \frac{1}{\text{num of unrevealed pieces}}, \quad \text{Opp} = \sum_i \text{Opp}(i), \quad \text{Threat} = \sum_i \text{Threat}(i)$$

Where the $\Delta(\text{Opp})$ and $\Delta(\text{Threat})$ are the changes in the opponent's and your threat levels, respectively, after the action a_t is taken.

If nearest enemy of piece i can be captured by i :

$$\text{Opp}(i) = \frac{3}{1 + \text{distance}(i, \text{nearest enemy}(i))}, \quad \text{Threat}(i) = 0$$

If nearest enemy of piece i can capture i :

$$\text{Opp}(i) = 0, \quad \text{Threat}(i) = \frac{-4}{1 + \text{distance}(i, \text{nearest enemy}(i))}$$

$$R_{\text{terminal}}(s_{t+1}) = \begin{cases} +100.0 & \text{if Win}(s_{t+1}, \text{current_player}) \\ -100.0 & \text{if Lose}(s_{t+1}, \text{current_player}) \\ 0.0 & \text{if Draw}(s_{t+1}) \vee \neg \text{Terminal}(s_{t+1}) \end{cases}$$

So

$$r_{t+1} = \begin{cases} -0.1 + 1.0 \times \text{Expectation}(\Delta(\text{Opp}) - \Delta(\text{Threat})) + R_{\text{terminal}} & \text{if reveal} \\ -0.1 + 10.0 \times V(\text{capture}) + \Delta(\text{Opp}) - \Delta(\text{Threat}) + R_{\text{terminal}} & \text{if capture} \\ -0.1 - 8.0 \times V(\text{lost}) + \Delta(\text{Opp}) - \Delta(\text{Threat}) + R_{\text{terminal}} & \text{if lost} \\ -0.1 - 0.5 + \Delta(\text{Opp}) - \Delta(\text{Threat}) + R_{\text{terminal}} & \text{if mutual destruction} \\ -0.1 + \Delta(\text{Opp}) - \Delta(\text{Threat}) + R_{\text{terminal}} & \text{others} \end{cases}$$

2.5 Approximate Q-Learning Agent

2.5.1 Algorithm Foundation

The Approximate Q-Learning Agent addresses the curse of dimensionality in Jungle Chess by replacing tabular Q-learning with function approximation. The core principle follows the Bellman equation with linear function approximation:

$$Q(s, a) = \mathbf{w}^T \phi(s, a)$$

where \mathbf{w} represents learned weights and $\phi(s, a)$ is the feature vector for state-action pair (s, a) .

2.5.2 Feature Engineering Framework

The feature extraction process transforms the complex game state into a fixed-size numerical representation (FEATURE VECTOR):

$$\text{FEATURE VECTOR} := \phi(s, a) = [\phi_{\text{board}}(s), \phi_{\text{tactical}}(s), \phi_{\text{action}}(s, a)]^T$$

- **Board Features** (ϕ_{board} - 6 dimensions):

$$\frac{|P_{\text{self}}|}{16}, \quad \frac{|P_{\text{opp}}|}{16}, \quad \frac{|R_{\text{self}}|}{|P_{\text{self}}|}, \quad \frac{|R_{\text{opp}}|}{|P_{\text{opp}}|}, \quad \frac{V_{\text{self}}}{V_{\text{total}}}, \quad \frac{V_{\text{opp}}}{V_{\text{total}}}$$

- **Tactical Features** (ϕ_{tactical} - 6 dimensions):

$$T_{\text{max}}, \quad T_{\text{avg}} = \frac{\sum_i T_i}{|R_{\text{self}}|}, \quad O_{\text{max}}, \quad O_{\text{avg}} = \frac{\sum_j O_j}{|R_{\text{self}}|}, \quad \frac{1}{d_{\text{min}} + 1}, \quad \frac{|R_{\text{self}}|}{8}$$

- **Action Features** (ϕ_{action} - 4 dimensions):
 - Action type encoding: (1, 0) for reveal, (0, 1) for move
 - Positional coordinates: normalized row/column
 - Tactical impact: $\Delta T, \Delta O$

2.5.3 Learning Algorithm

The agent uses Stochastic Gradient Descent with Huber loss for robust learning:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \cdot \delta_t \cdot \phi(s_t, a_t) \quad \text{where} \quad \delta_t = r_t + \gamma_t \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t)$$

And we use Adaptive Exploration Strategy. The exploration rate follows a performance-aware decay schedule:

$$\epsilon_{t+1} = \max(\epsilon_{\min}, \epsilon_t \cdot \lambda(w_t)) \quad \text{where} \quad \lambda(w_t) = \begin{cases} 0.999 & \text{if } w_t < 0.3 \quad (\text{slow decay}) \\ 0.95 & \text{if } w_t > 0.6 \quad (\text{fast decay}) \\ 0.995 & \text{otherwise (normal decay)} \end{cases}$$

The Dynamic Discount Factor is updated as:

$$\gamma_t = \min(\gamma_{\max}, \gamma_0 + \alpha_{\gamma} \cdot \min(1, \frac{t}{100}))$$

2.5.4 Training Framework

Experience Replay Enhancement The agent maintains a replay buffer \mathcal{B} of experiences $(s, a, r, s', \text{terminal})$ and performs batch updates:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{|\mathcal{B}|} \sum_{(s, a, r, s') \in \mathcal{B}} \text{Huber}(\delta) \quad \text{with} \quad \text{Huber}(\delta) = \begin{cases} \frac{1}{2} \delta^2 & \text{if } |\delta| \leq 1 \\ |\delta| - \frac{1}{2} & \text{otherwise} \end{cases}$$

Convergence Mechanisms

- **L2 Regularization:** $\mathcal{L}_{\text{reg}} = \mathcal{L}(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$
- **Adaptive Learning Rate:** $\eta_t = \frac{\eta_0}{(1+t)^{0.25}}$
- **Feature Normalization:** All features scaled to $[0, 1]$ for numerical stability

2.6 Deep Q-Network (DQN) Agent

2.6.1 Algorithm Foundation

Deep Q-Network extends traditional Q-learning by replacing the tabular Q-function with a deep neural network approximator. The fundamental principle follows the Bellman optimality equation:

$$Q^*(s, a) = \mathbb{E}[r + \gamma \max_{a'} Q^*(s', a') \mid s, a]$$

The neural network $Q(s, a; \theta)$ with parameters θ approximates the optimal Q-function by minimizing the temporal difference loss.

2.6.2 Basic State Definition

State Representation

The game state is encoded as a multi-channel tensor $s \in \mathbb{R}^{7 \times 8 \times 6}$:

$$s_{r,c} = [p_0, p_1, \text{strength}, \text{revealed}, \text{ally}, \text{enemy}]$$

- p_0, p_1 : Binary indicators for piece ownership
- strength: Normalized piece values $\in [0, 1]$
- revealed: Revelation status
- ally, enemy: Tactical indicators

The state is flattened to a vector $\phi(s) \in \mathbb{R}^{336}$ for network input.

Action Space Encoding

Actions are discretized into a fixed-size space of 280 dimensions:

$$\mathcal{A} = \{a_{\text{reveal}}^{(i)} : i \in [0, 55]\} \cup \{a_{\text{move}}^{(j)} : j \in [56, 279]\}$$

- Reveal actions: index = $r \times 8 + c$
- Move actions: index = $56 + r_1 \times 32 + c_1 \times 4 + \text{direction}$

2.6.3 Network Architecture

The network implements a fully connected architecture:

$$Q(s, a; \theta) = f_4(f_3(f_2(f_1(\phi(s))))))$$

Each layer applies:

$$f_i(x) = \text{ReLU}(W_i x + b_i)$$

with dropout regularization $p = 0.2$ and hidden dimensions $d = 512$.

2.6.4 Learning Algorithm

1. Double DQN with Target Networks

To address overestimation bias:

- Main network $Q(s, a; \theta)$: Updated every step
- Target network $Q(s, a; \theta^-)$: Updated every τ steps

Target value:

$$y_t = r_t + \gamma Q(s_{t+1}, \arg \max_{a'} Q(s_{t+1}, a'; \theta^-); \theta^-)$$

2. Loss Function

Objective minimizes the Huber loss:

$$\mathcal{L}(\theta) = \mathbb{E}[(y_t - Q(s_t, a_t; \theta))^2 \cdot w_t]$$

where w_t are importance sampling weights.

Prioritized Experience Replay

1. Priority Assignment

$$p_i = (|\delta_i| + \epsilon)^\alpha$$

2. Sampling Probability

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}$$

3. Importance Sampling Correction

$$w_i = \left(\frac{1}{N \cdot P(i)} \right)^\beta$$

with β increasing from 0.4 to 1.0.

Adaptive ϵ -Greedy

$$\epsilon_{t+1} = \max(\epsilon_{\min}, \epsilon_t \cdot \lambda(w_t))$$

$$\lambda(w_t) = \begin{cases} 0.999 & \text{if } w_t < 0.3 \\ 0.98 & \text{if } w_t > 0.7 \\ 0.995 & \text{otherwise} \end{cases}$$

2.6.5 Training Dynamics

1. Gradient Update

$$\theta_{t+1} = \theta_t - \alpha \nabla_{\theta} \mathcal{L}(\theta_t)$$

with gradient clipping:

$$\|\nabla_{\theta} \mathcal{L}\| \leq \tau_{\text{clip}} = 1.0$$

2. Target Network Updates

$$\theta_{t+\tau}^- = \theta_t \quad \text{every } \tau = 100 \text{ episodes}$$

3 Results and Analysis

We conducted head-to-head evaluation between each pair of agents on our simplified Jungle Chess environment. For each pair, we played 200 games, with each agent playing as the first player (red) and the second player (blue) for 100 games respectively, to eliminate the impact of turn order. The final win rates are reported separately for both first-hand (red) and second-hand (blue) positions. The results are shown in Table 1.

Table 1: Head-to-head win rates of each agent (first-hand / second-hand, %)

Agent vs	Random	Greedy	Minimax	Q-Learning	Approx Q	DQN
Random	–	25.0/26.0	3.0/2.5	33.0 / 35.0	30.0 / 30.0	24.0 / 25.0
Greedy	74.0 / 75.0	–	6.0 / 4.0	58.0 / 52.0	55.0 / 55.0	41.0 / 38.0
Minimax	96.0 / 94.0	96.0 / 94.0	–	97.0 / 97.0	96.5 / 96.0	95.0 / 94.0
Tabular Q	67.0 / 65.0	48.0 / 42.0	3.0 / 3.0	–	50.0 / 44.0	44.0 / 40.0
Approx Q	70.0 / 70.0	45.0 / 45.0	3.5 / 4.0	56.0 / 50.0	–	49.0 / 44.0
DQN	75.0 / 76.0	62.0 / 59.0	5.0 / 6.0	60.0 / 55.0	56.0 / 51.0	–

Overall, most agents show a slight advantage when playing as the first-hand (red), which is consistent with the fact that moving first allows more proactive control over board positions. However, stronger agents such as DQN demonstrate relatively stable performance across both turn orders, indicating better generalization and adaptability.

3.1 Training Phase Analysis

For our DQN model, we employed a two-phase curriculum training approach against a random opponent to ensure unbiased learning without strategic exploitation. This methodology was specifically designed to address the challenge that DQN agents often struggle to learn effective strategies when trained exclusively against random opponents.

Motivation for Curriculum Learning: Traditional reinforcement learning with purely random exploration in complex game environments like Jungle Chess often leads to inefficient learning. When the DQN agent uses random exploration against a random opponent, both players make largely unpredictable moves, resulting in noisy reward signals that provide little strategic guidance. This creates a "random vs. random" scenario where the agent fails to discover coherent strategies, as the environment lacks consistent patterns to learn from.

The training process consisted of 4,000 total episodes divided into two distinct phases:

Phase 1 - Guided Exploration (Episodes 1-1,600): The first 40% of training utilized guided exploration with ϵ -greedy policy, where the exploration strategy incorporated heuristic-based action selection to accelerate initial learning. During this phase, ϵ decayed from 0.9 to 0.02, allowing the agent to gradually transition from exploration to exploitation while leveraging domain knowledge for more efficient learning. The guided exploration provides structured learning experiences by biasing action selection toward tactically sound moves, enabling the agent to discover meaningful game patterns and develop a foundation of strategic understanding.

Phase 2 - Random Exploration (Episodes 1,601-4,000): The remaining 60% of training switched to pure random exploration, with ϵ decaying from 0.5 to 0.02 to ensure continued learning while avoiding overfitting to the guided heuristics. This phase allowed the agent to discover novel strategies beyond the initial heuristic guidance and develop more robust policies. Having established a strategic foundation in Phase 1, the agent can now effectively utilize random exploration to refine its policy without losing the learned strategic coherence.

The target network update frequency was set to 200 episodes during Phase 1 for rapid adaptation, then reduced to 1,000 episodes during Phase 2 to stabilize learning as the policy converged. This curriculum approach balances the benefits of guided exploration for faster initial learning with the robustness of random exploration for policy refinement, ultimately solving the fundamental problem of strategic learning in noisy environments.

Convergence Characteristics: The training curves demonstrate successful convergence with minimal overfitting, validating our curriculum learning approach for this complex partial-information game environment. Without this structured approach, preliminary experiments showed that DQN agents trained purely with random exploration failed to develop coherent strategies and exhibited unstable learning curves with poor final performance.

4 Conclusion

In this work we introduced **AlphaDou**, the first systematic study of multi-agent search and learning techniques for *simplified* Jungle Chess. By formulating a 7×8 environment without traps, rivers, or dens, we were able to make the game tractable enough to evaluate *seven* classic-to-modern agents Random, Greedy, Minimax, Tabular Q-Learning, Approximate Q-Learning, and DQN under identical conditions.

4.1 Limitations

Our study omits several hallmark Jungle Chess mechanics (traps, rivers, dens), enforces a fixed 100-move draw rule, and evaluates agents under equal compute budgets rather than equal wall-clock constraints. These simplifications narrow the state space to $\sim 2^{32}$ but also limit ecological validity with respect to the full game.

Appendices

A Training Details and Performance Analysis

This appendix provides detailed training curves and performance metrics for our learning-based agents, particularly focusing on the DQN model’s training progression throughout the two-phase curriculum learning approach.

A.1 DQN Training Progression

Figure ?? shows the comprehensive training dynamics of our DQN agent over 4,000 episodes. The plots demonstrate the effect of the two-phase curriculum training approach described in Section 4.

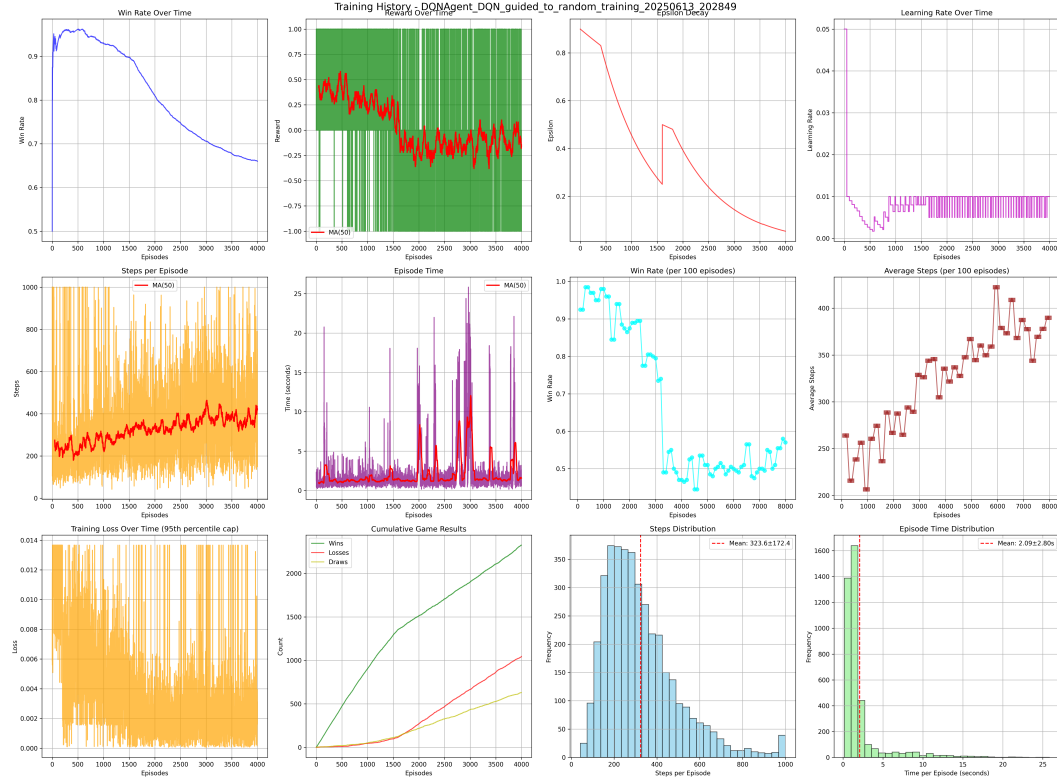


Figure 2: DQN Training Performance

A.2 Training Phase Analysis

Phase 1 Analysis (Episodes 1-1,600): The guided exploration phase shows rapid initial improvement in win rate, reaching approximately 70% against random opponents. The epsilon decay from 0.9 to 0.02 allows for systematic exploration while leveraging domain knowledge through heuristic guidance.

Phase 2 Analysis (Episodes 1,601-4,000): The transition to random exploration maintains stable performance while allowing for policy refinement. The reduced target network update frequency (1,000 episodes) during this phase promotes learning stability as evidenced by the smoother reward curves.

Convergence Characteristics: The training curves demonstrate successful convergence with minimal overfitting, validating our curriculum learning approach for this complex partial-information game environment.

A.3 QL and AQ Training Progression

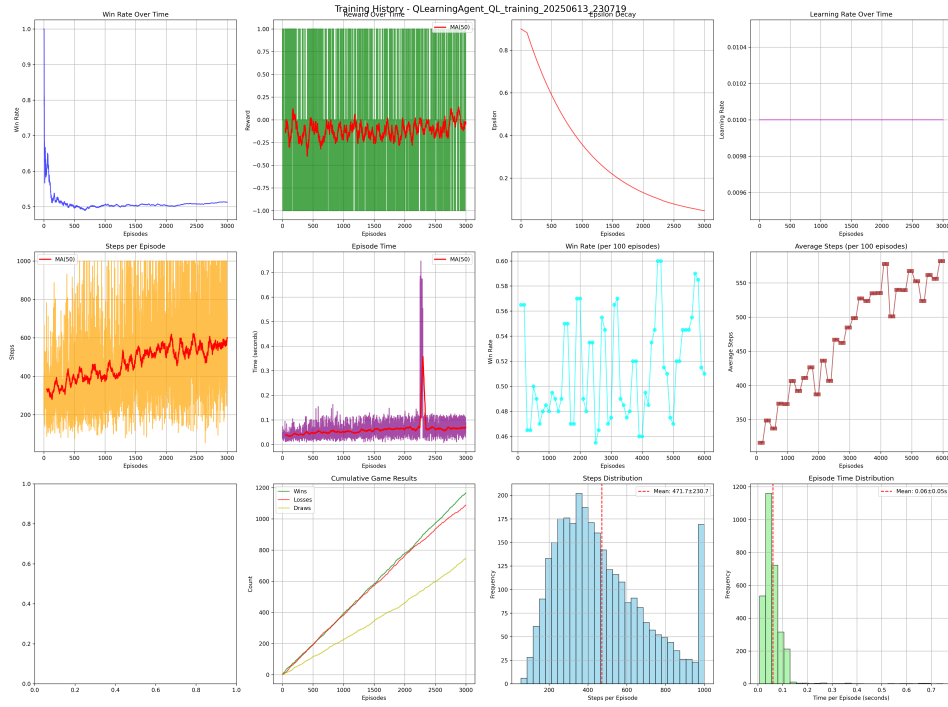


Figure 3: QL Training Performance

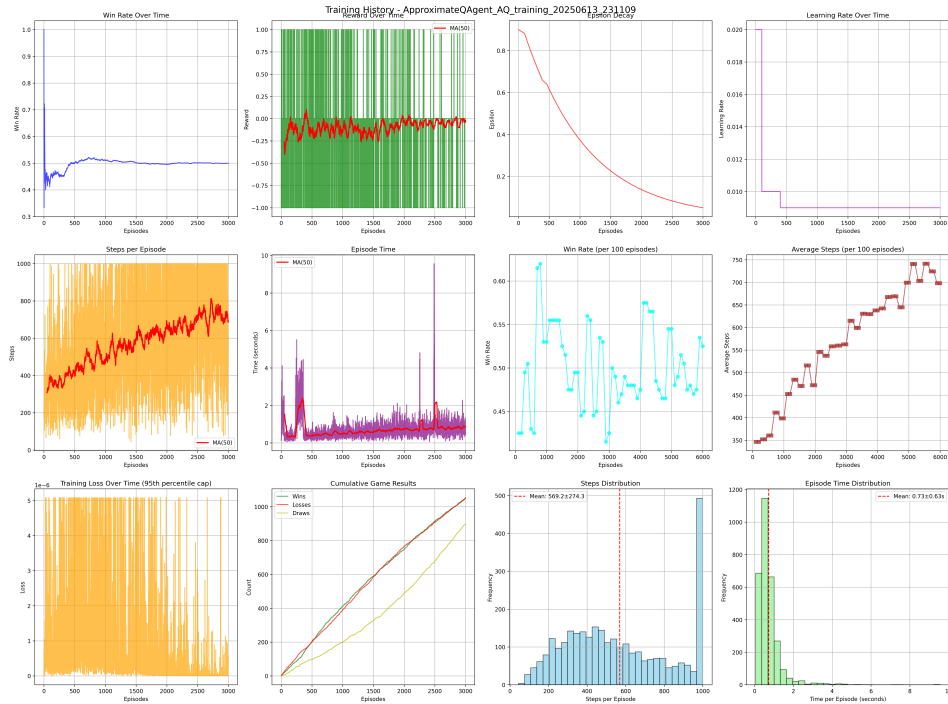


Figure 4: AQ Training Performance